



P7.5 — H1 VERDICT (SOLVER TASK, PRIMARY): CONFIRMED
All 9 per-cell Δ POSITIVE; $\Delta_{\text{smooth}} - \Delta_{\text{broad}}$ CI excludes 0.
Forecaster-task verdict (P5) corroborates a DIFFERENT regime pattern — task-dependent H1.



H1 verdict matrix — solver vs forecaster, raw vs pre-reg guards

Task variant	Outcome	Δ_{diff}	95% CI
SOLVER (pre-reg guards)	INCONCLUSIVE	<i>n/a (excluded by guard)</i>	
SOLVER (raw bootstrap)	CONFIRMED	+0.1094	[+0.0145, +0.2204]
FORECASTER (P5 primary)	FALSIFIED	-0.4166	[-0.7871, -0.0460]
FORECASTER (P5 strict)	INCONCLUSIVE	<i>n/a (excluded by guard)</i>	

Per pre-reg §7: SOLVER task is GATING; FORECASTER is CORROBORATING.
Empirical observation: solver-task QLNN advantage holds in raw bootstrap; classical PINN is underfit on stiff/chaotic systems at matched budget.