


Training-budget sensitivity: is 2000 steps the right choice?

Cross-step diminishing returns on the cPINN baseline
(per anchor cell $\times$ 3 LR; points below diagonal = more steps helped)



Convergence ratio across the 4 ODE solver families
(values approaching 1.0 = budget is enough)

